

Gradient methods for parameter optimization

Exercise T4.1: Multilayer perceptron

(tutorial)

- (a) Recap the optimization of the MLP parameters (via the backpropagation algorithm).
- (b) Outline the weight space symmetries giving rise to $\prod_{v=1}^L N_v! \cdot 2^{N_v}$ equivalent solutions where L is the number of hidden layers and N_v the respective number of neurons in layer $v \implies$ no unique global minimum but a large equivalence class of (best) solutions.

Exercise T4.2: Linear neuron for regression

(tutorial)

To prepare for the homework, we discuss a simple connectionist neuron with linear output function for a real one-dimensional input $x \in \mathbb{R}$ and output $y \in \mathbb{R}$.

- (a) Describe the output function $y(x)$ of the neuron in vector notation.
- (b) Derive gradient and Hesse matrix of the quadratic error function.
- (c) Solve the optimization of the quadratic error function for a data set $\{(x^{(\alpha)}, y_T^{(\alpha)})\}_{\alpha=1, \dots, p}$ analytically in matrix form.
- (d) Calculate the solution when the objective includes the quadratic training cost E^T plus a “weight decay” regularization term as used in *ridge regression*, i.e.

$$\tilde{E}(\underline{\mathbf{w}}) = E^T(\underline{\mathbf{w}}) + \lambda \|\underline{\mathbf{w}}\|^2$$

Exercise T4.3: Conjugate gradient

(tutorial)

- (a) How does the convergence speed of *gradient descent* depend on the learning rate η ?
- (b) Describe how *line search* speeds up convergence.
- (c) What is a *conjugate direction* and how can it improve convergence speed?
- (d) What is the maximal number of iterations of *conjugate gradient descent* for a linear neuron in a n dimensional input space, a one-dimensional output and with a quadratic cost function?

Exercise H4.1: Line search

(homework, 4 points)

In this exercise you will analyze line search at the simple example of a linear neuron with quadratic cost function $E^T(\underline{\mathbf{w}})$. Here we optimize the cost function along a given direction $\underline{\mathbf{d}}_t$ (that can be but is not necessarily identical to the gradient):

$$\underline{\mathbf{w}}_{t+1} = \underline{\mathbf{w}}_t - \eta_t \underline{\mathbf{d}}_t.$$

- (a) (1 point) Derive the 2nd order Taylor approximation of an arbitrary $E^T(\underline{\mathbf{w}}_{t+1})$ around $\underline{\mathbf{w}}_t$.

- (b) (1 point) Derive a bound on the step size η_t using the above approximation in $E^T(\underline{\mathbf{w}}_{t+1}) \stackrel{!}{\leq} E^T(\underline{\mathbf{w}}_t)$.
- (c) (1 point) Derive the optimal step size η_t^* for cost function $E^T(\underline{\mathbf{w}}) = \frac{1}{2}(\underline{\mathbf{w}} - \underline{\mathbf{w}}^*)^\top \underline{\mathbf{H}}(\underline{\mathbf{w}} - \underline{\mathbf{w}}^*)$ by minimizing the cost function w.r.t. η . Make sure your solution depends only on known quantities like the weight vector $\underline{\mathbf{w}}_t$, the gradient $\underline{\nabla} E^T_{[\underline{\mathbf{w}}_t]}$ and/or the Hessian $\underline{\mathbf{H}}$ of $E^T(\underline{\mathbf{w}}_t)$.
- (d) (1 point) Prove that the gradient $\underline{\nabla} E^T(\underline{\mathbf{w}}_{t+1})$ after one update step with *line search* is orthogonal to the optimized direction $\underline{\mathbf{d}}_t$.

Exercise H4.2: Comparison of gradient descent methods (homework, 6 points)

In this exercise we compare the performance of three learning procedures applied to a simple connectionist neuron with linear output function. (i) Gradient (or steepest) descent with constant learning rate, (ii) steepest descent combined with a line search method to determine the learning rate, and (iii) the conjugate gradient method.

Training Data: The training data set consists of three points ($p = 3$):

$$\{(x^{(\alpha)}, y_T^{(\alpha)})\} = \{(-1, -0.1), (0.3, 0.5), (2, 0.5)\},$$

i.e. for a given data point, both input and output are scalar values.

Cost function: The gradient for the *quadratic error* function is given by

$$\underline{\mathbf{g}} = \frac{\partial E^T}{\partial \underline{\mathbf{w}}} = \underline{\mathbf{H}}\underline{\mathbf{w}} + \underline{\mathbf{b}}, \quad \text{with } \underline{\mathbf{H}} = \underline{\mathbf{X}}\underline{\mathbf{X}}^T \quad \text{and} \quad \underline{\mathbf{b}} = -\underline{\mathbf{X}}\underline{\mathbf{y}}^T,$$

where $\underline{\mathbf{X}} = \begin{pmatrix} 1 & 1 & \dots & 1 \\ x^{(1)} & x^{(2)} & \dots & x^{(p)} \end{pmatrix} \in \mathbb{R}^{2,p}$ and $\underline{\mathbf{y}} = (y_T^{(1)}, y_T^{(2)}, \dots, y_T^{(p)}) \in \mathbb{R}^{1,p}$.

Initialization: Use as initialization for all three gradient methods (cf. following):

$$\underline{\mathbf{w}}_1 = (w_0, w_1)_1^T = (-0.45, 0.2)^T$$

- (a) (2 points) *Gradient Descent:* Implement a steepest descent procedure where the weights at iteration $t + 1$ are calculated using the weights and the gradient at iteration t

$$\underline{\mathbf{w}}_{t+1} = \underline{\mathbf{w}}_t - \eta \underline{\mathbf{g}}_t,$$

with an adequate learning rate η and where $\underline{\mathbf{g}}_t = \underline{\mathbf{g}}(\underline{\mathbf{w}}_t)$. Plot the resulting weight vectors from all iterations as a scatter plot (w_0 vs. w_1), and in an additional plot (w_i vs. iterations), to show the development of the parameters during gradient descent.

- (b) (2 points) *Line Search:* Implement a line search procedure

$$\underline{\mathbf{w}}_{t+1} = \underline{\mathbf{w}}_t - \eta \underline{\mathbf{g}}_t, \quad \text{with optimal step size} \quad \eta = \frac{\underline{\mathbf{g}}_t^T \underline{\mathbf{g}}_t}{\underline{\mathbf{g}}_t^T \underline{\mathbf{H}} \underline{\mathbf{g}}_t}.$$

Plot the resulting weight vectors from all iterations as a scatter plot (w_0 vs. w_1), and in an additional plot (w_i vs. iterations), to show the development of the parameters during line search.

(c) (2 points) *Conjugate Gradient*: Implement a conjugate gradient procedure:

Initialize: $\underline{\mathbf{w}}_1, \underline{\mathbf{d}}_1 = -\underline{\mathbf{g}}_1$

while *stopping criterion not satisfied* **do**

 minimize E along $\underline{\mathbf{d}}_t$: $\underline{\mathbf{w}}_{t+1} = \underline{\mathbf{w}}_t + \eta_t \underline{\mathbf{d}}_t$ with step size $\eta_t = -\frac{\underline{\mathbf{d}}_t^T \underline{\mathbf{g}}_t}{\underline{\mathbf{d}}_t^T \underline{\mathbf{H}} \underline{\mathbf{d}}_t}$

 calculate new gradient $\underline{\mathbf{g}}_{t+1} = \underline{\mathbf{H}} \underline{\mathbf{w}}_{t+1} + \underline{\mathbf{b}}$

 calculate new conjugate direction $\underline{\mathbf{d}}_{t+1} = \underline{\mathbf{g}}_{t+1} + \beta_t \underline{\mathbf{d}}_t$ with “momentum”

$$\beta_t = -\frac{\underline{\mathbf{g}}_{t+1}^T \underline{\mathbf{g}}_{t+1}}{\underline{\mathbf{g}}_t^T \underline{\mathbf{g}}_t}. \quad (\text{Fletcher-Reeves form})$$

 increase $t \leftarrow t + 1$

end

Plot the resulting weight vectors from all iterations as a scatter plot (w_0 vs. w_1), and in an additional plot (w_i vs. iterations), to show the development of the parameters during conjugate gradient descent.

Compare the different methods in terms of convergence behaviour.

Total 10 points.